

A vertical strip on the left side of the slide features an abstract background of glowing, diagonal lines in various colors like blue, red, yellow, and green, set against a dark background.

# Concurrent Programming

This presentation introduces concurrent programming. We will explore its concepts and applications. We will also examine the history and key differences. Let's delve into the world of concurrency.

# Defining Concurrency

## Parallel Execution

Tasks run simultaneously on multi-core systems.

## Interleaved Execution

Tasks appear to run at the same time.



# Why Study Concurrency?

- 1 Real-World Systems**  
It models complex systems effectively.
- 2 Various Architectures**  
It is useful across computer architectures.
- 3 Application Necessity**  
It's vital for certain applications.



# Concurrent Application Examples



Web servers handle multiple client requests.



Artificial Intelligence Agents



Data Science processing



Simulations forecast weather or simulate flights.



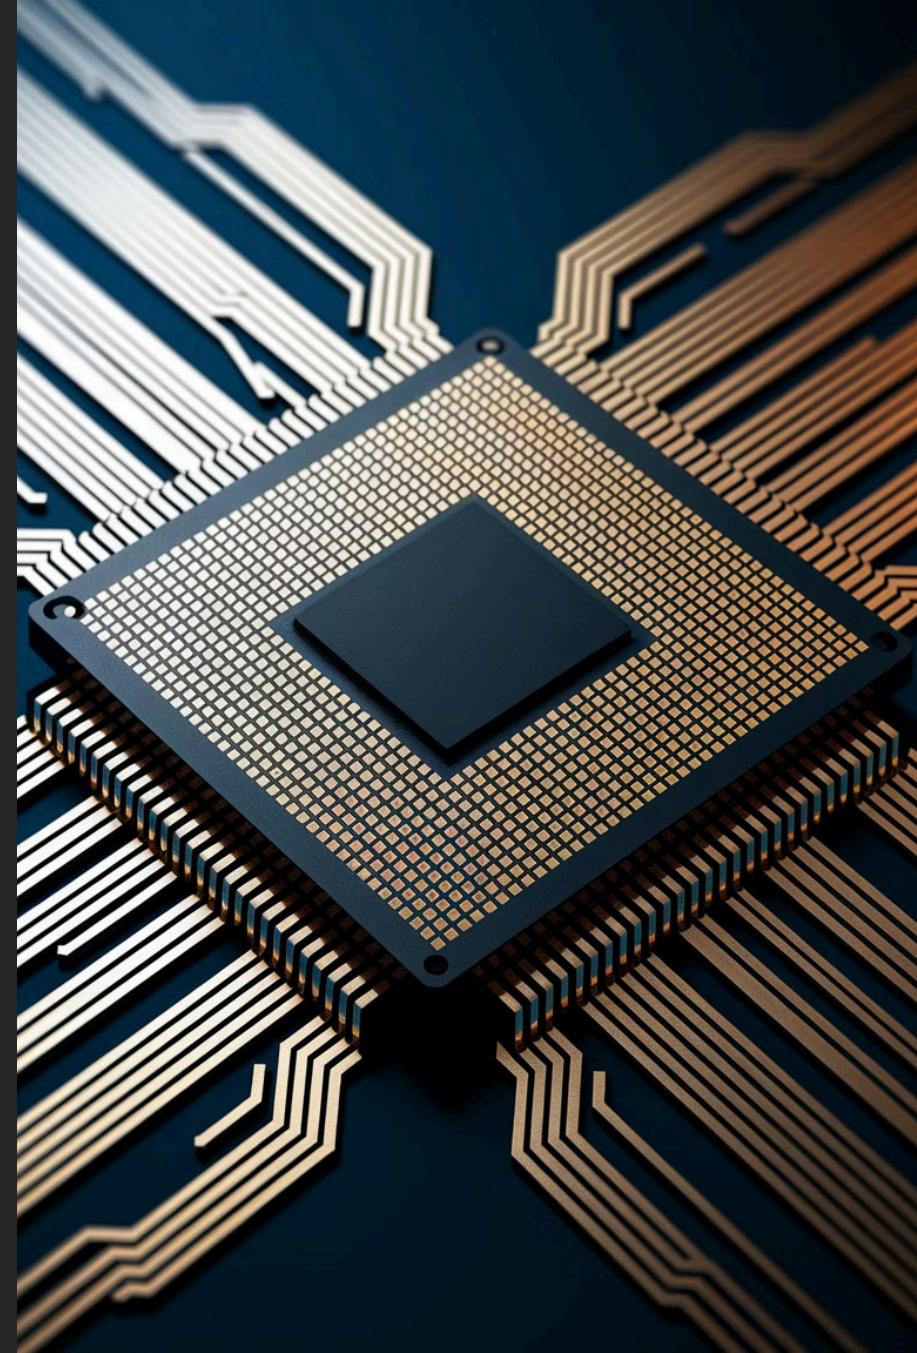
Games manage actions from multiple players.



Operating systems running multiple applications.

# Introduction to Moore's Law

Moore's Law predicts processing power growth. It has driven innovation for decades. Let's look at its origin and impact.





# Moore's Law: Definition

## Transistors

The number of transistors doubles every two years.

## Cost

The cost of computers is halved.

## Origin

Coined by Gordon Moore in 1965.



# Moore's Law: Impact

## 1 Exponential Growth

Computing power and efficiency increase.

## 2 Innovation Acceleration

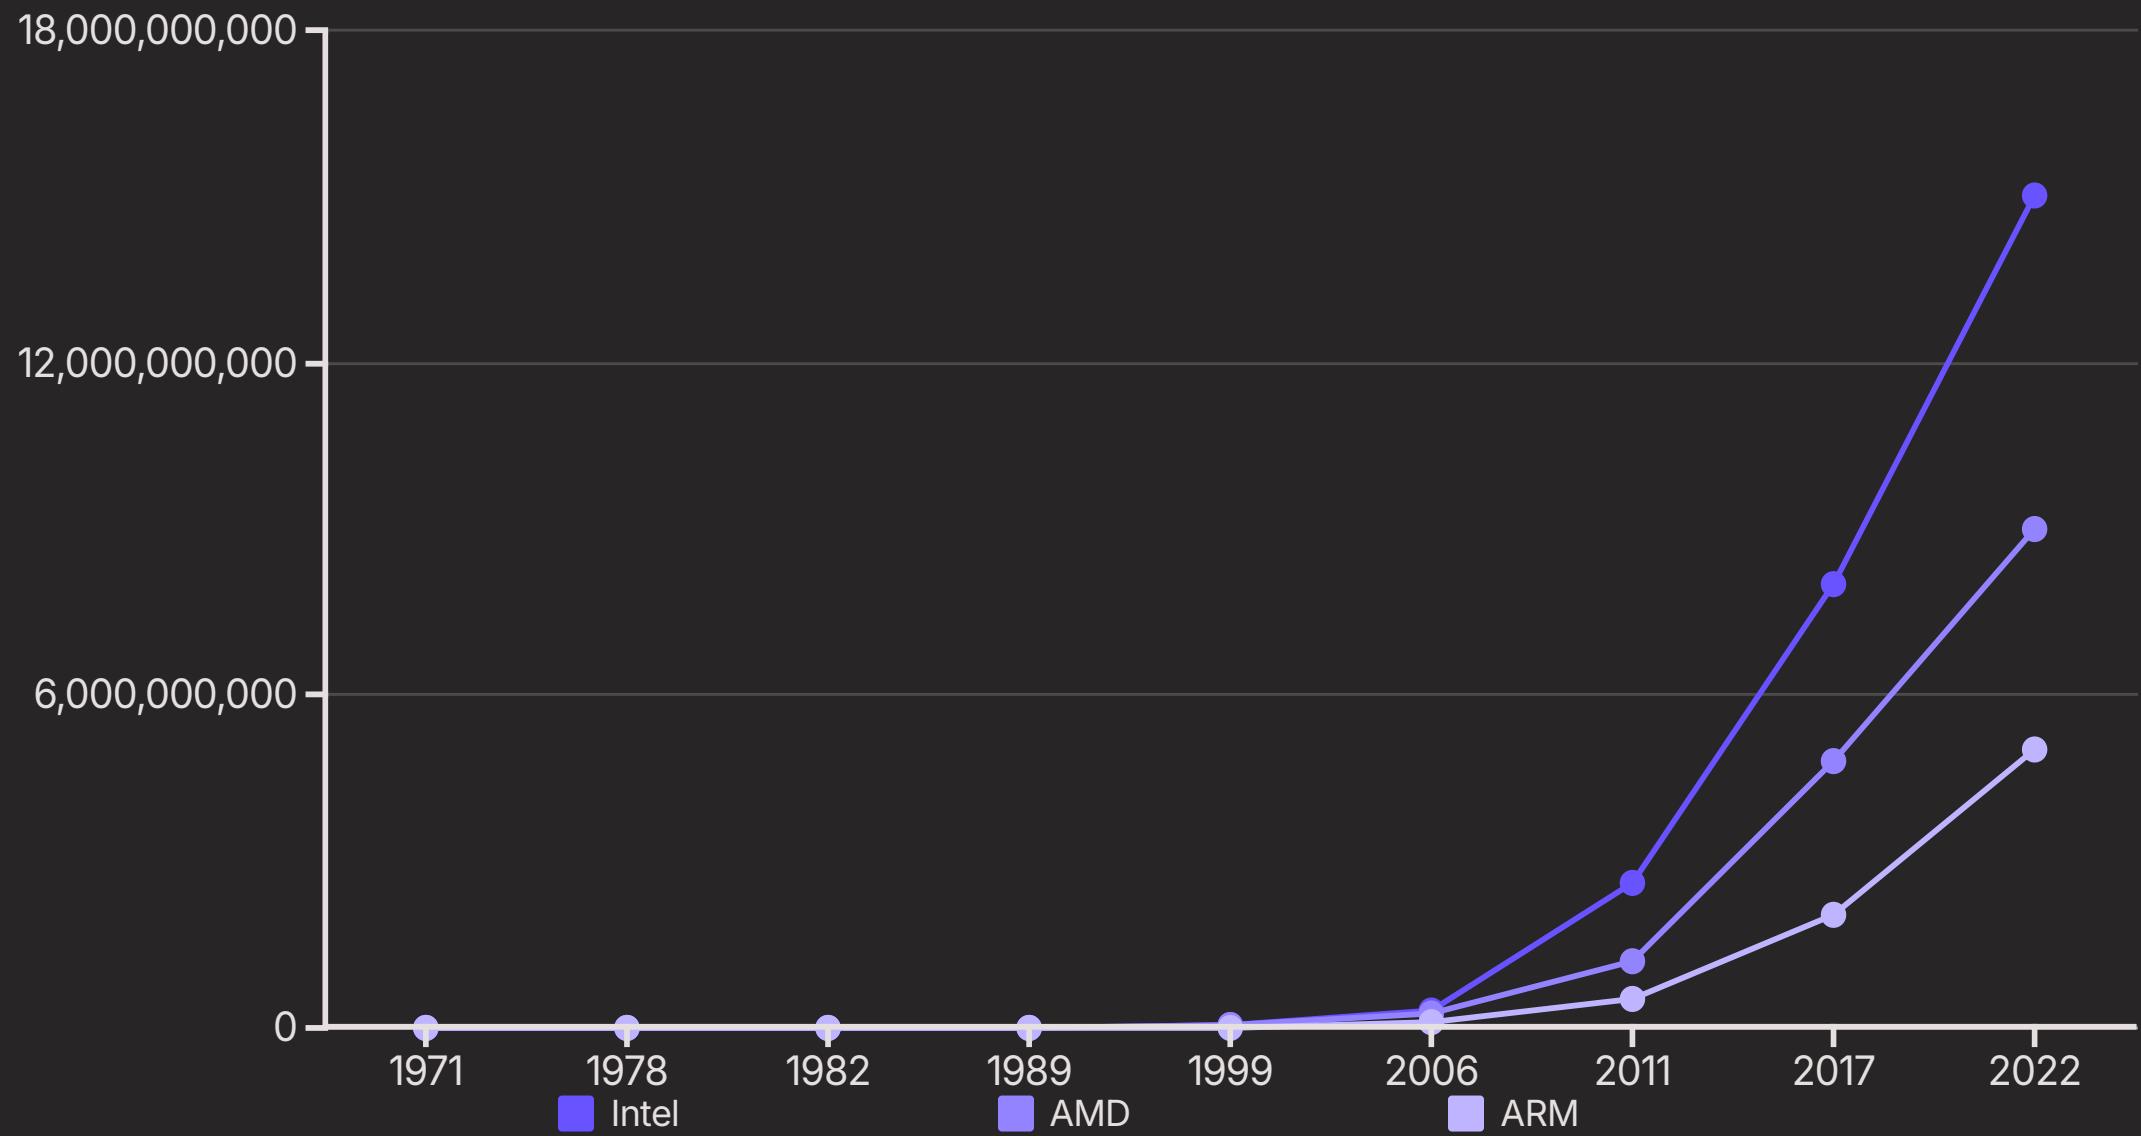
More powerful and compact devices emerge.

## 3 Economic Effects

The tech industry emphasizes rapid development.

# Transistors Over Time

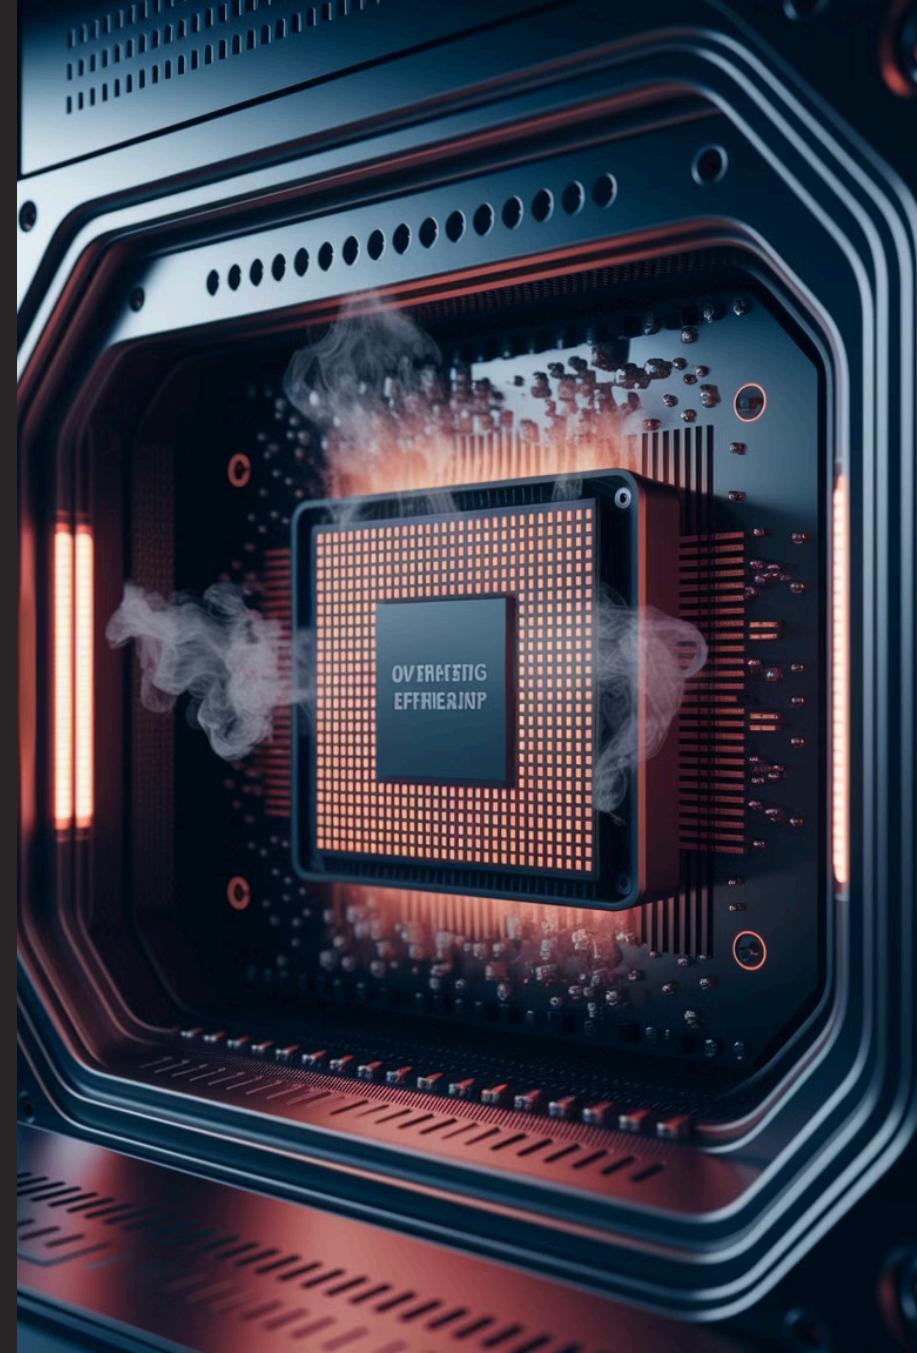
Transistor count has grown exponentially across processor generations, as predicted by Moore's Law.



Key milestones include: Intel 4004 (1971, 2,300 transistors), Intel 8086 (1978, 29,000), Intel 80386 (1982, 134,000), Intel Pentium (1993, 3.1 million), and modern processors with billions of transistors. This exponential growth has powered decades of computing advancement.

# Limits in Processor Speed

Speed improvements began to slow in the mid-2000s. The number of transistors continued to grow.



# Heat and Power Wall

## Heat Generation

Transistor density increases heat generation.

## Power Consumption

Higher clock speeds increase power use and heat.

## Diminishing Returns

Beyond certain point Clock speed gains do not translate to performance gains.

# Innovations Beyond Clock Speed

- 1
- 2
- 3

## Architectural Improvements

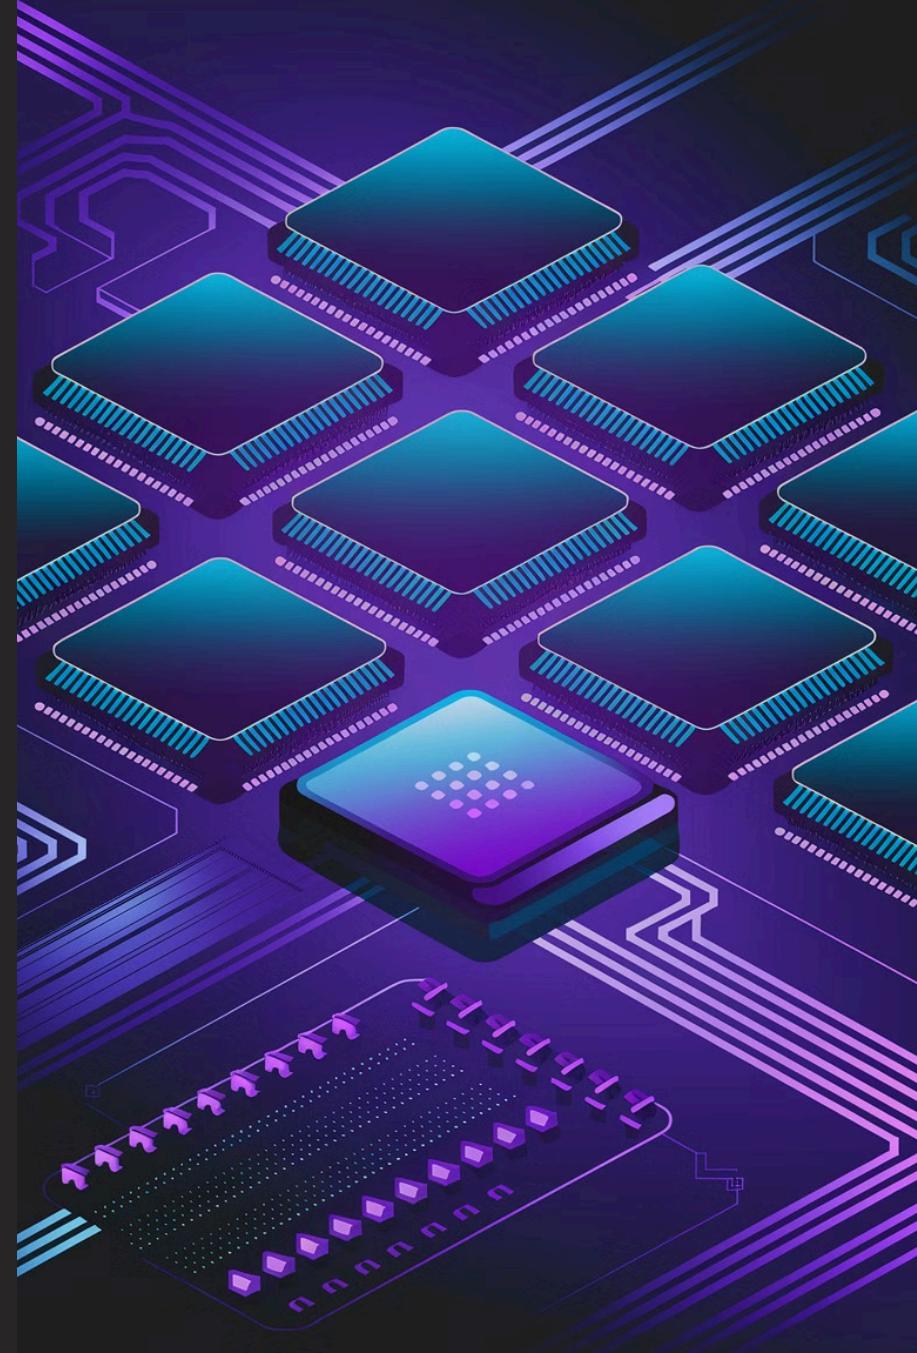
Advances in CPU architecture, larger caches.

## Multi-core Processors

Adding more processing cores for parallel execution.

## Specialized Units

Using specialized computing units, like GPUs.





# Parallelism and Hardware

Multicore Processors

Multiple Processors

GPUs

Grid Computing on Networks

# History of Concurrency

Let's look at the early days of computing. We will discuss multiprogramming and time-sharing. This will give us context.



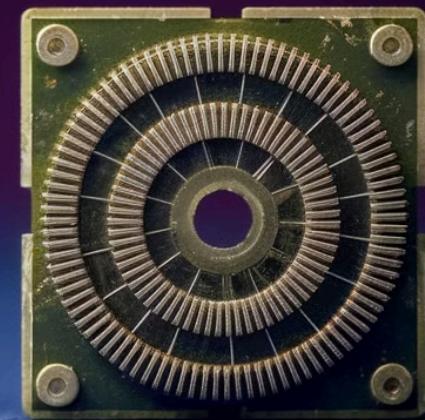


# Early Days of Computing

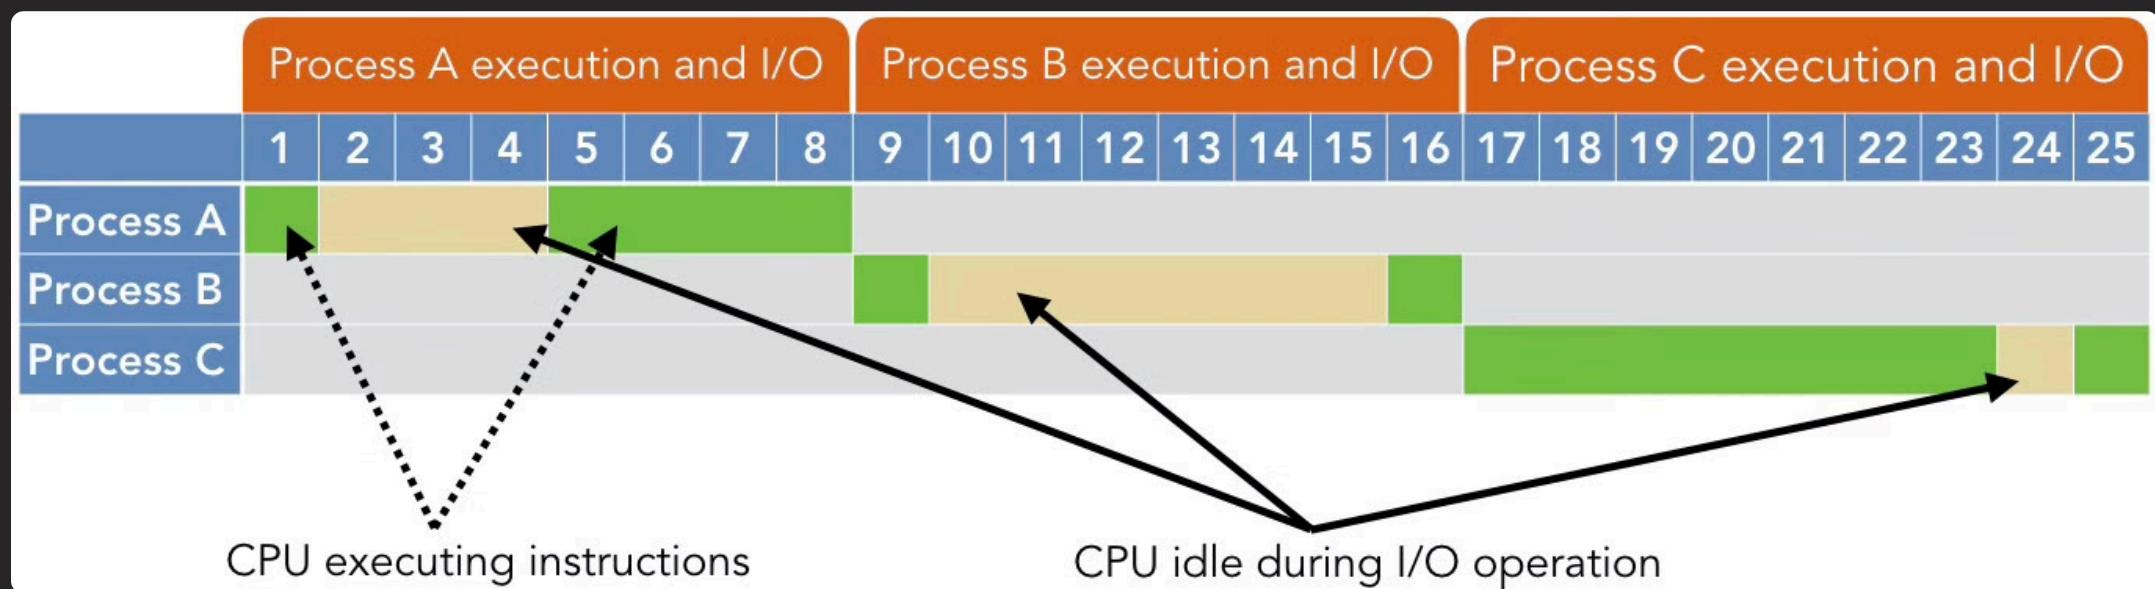
- 1 Computers ran a single program.
- 2 Direct access to all machine resources.
- 3 Inefficient resource utilization.

# Introduction to Multi-programming

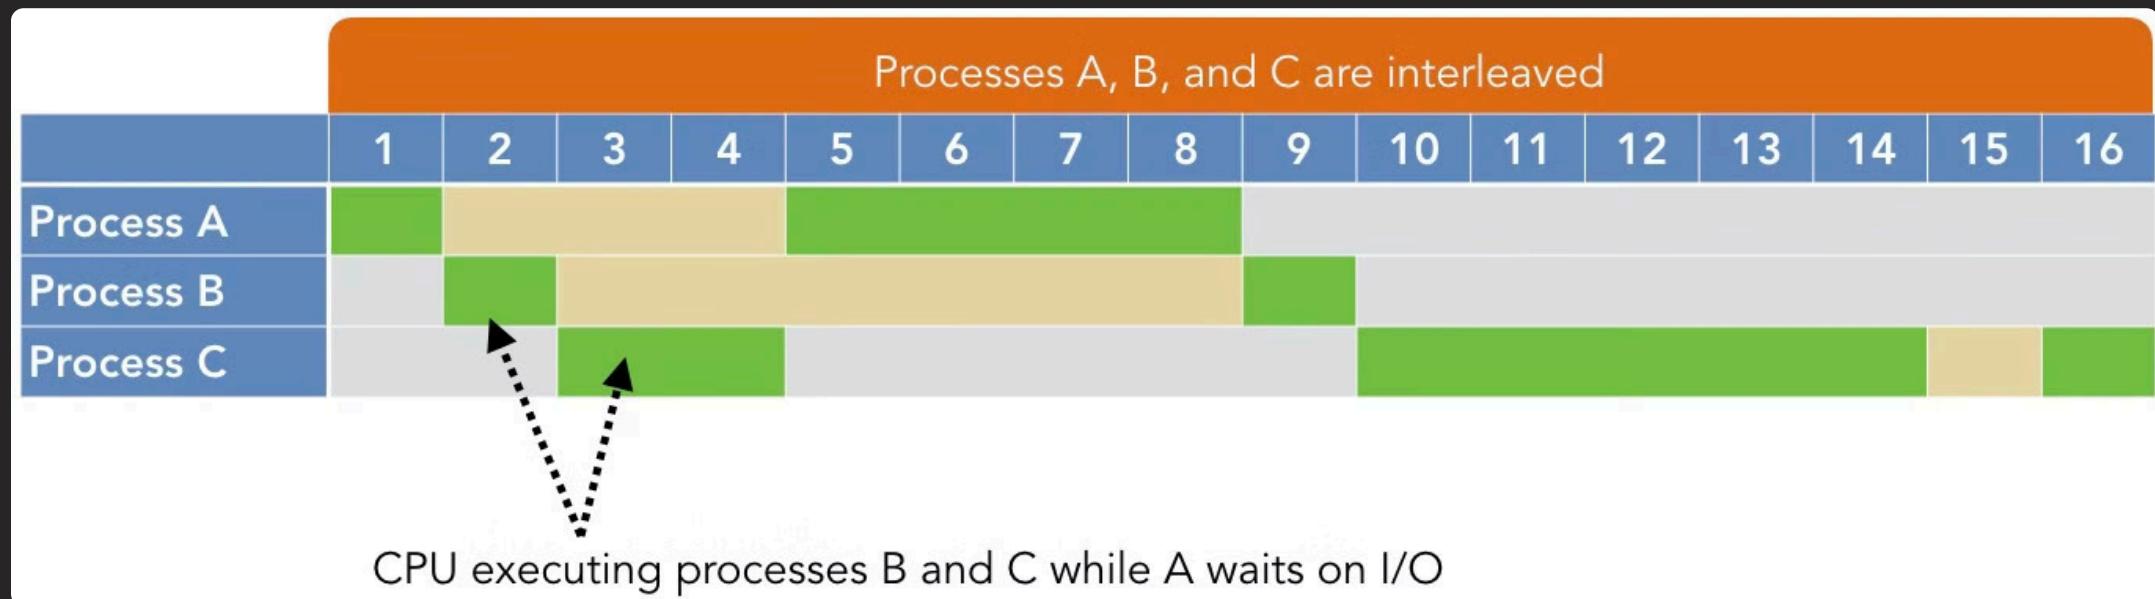
- 1** Shift from single to multiple program execution.
- 2** Loading several jobs into memory to be executed concurrently
- 3** Enhancing CPU utilization. by organizing jobs so CPU is always busy



# Uni-programming: sequential jobs



# Multiprogramming: parallel jobs



# Operating Systems and Time-Sharing

## Time-Sharing

Time-sharing is a logical extension of multiprogramming.

## Processor Sharing

Processor's time is shared among multiple users.

## User Experience

Creates the illusion of a dedicated machine for each user.





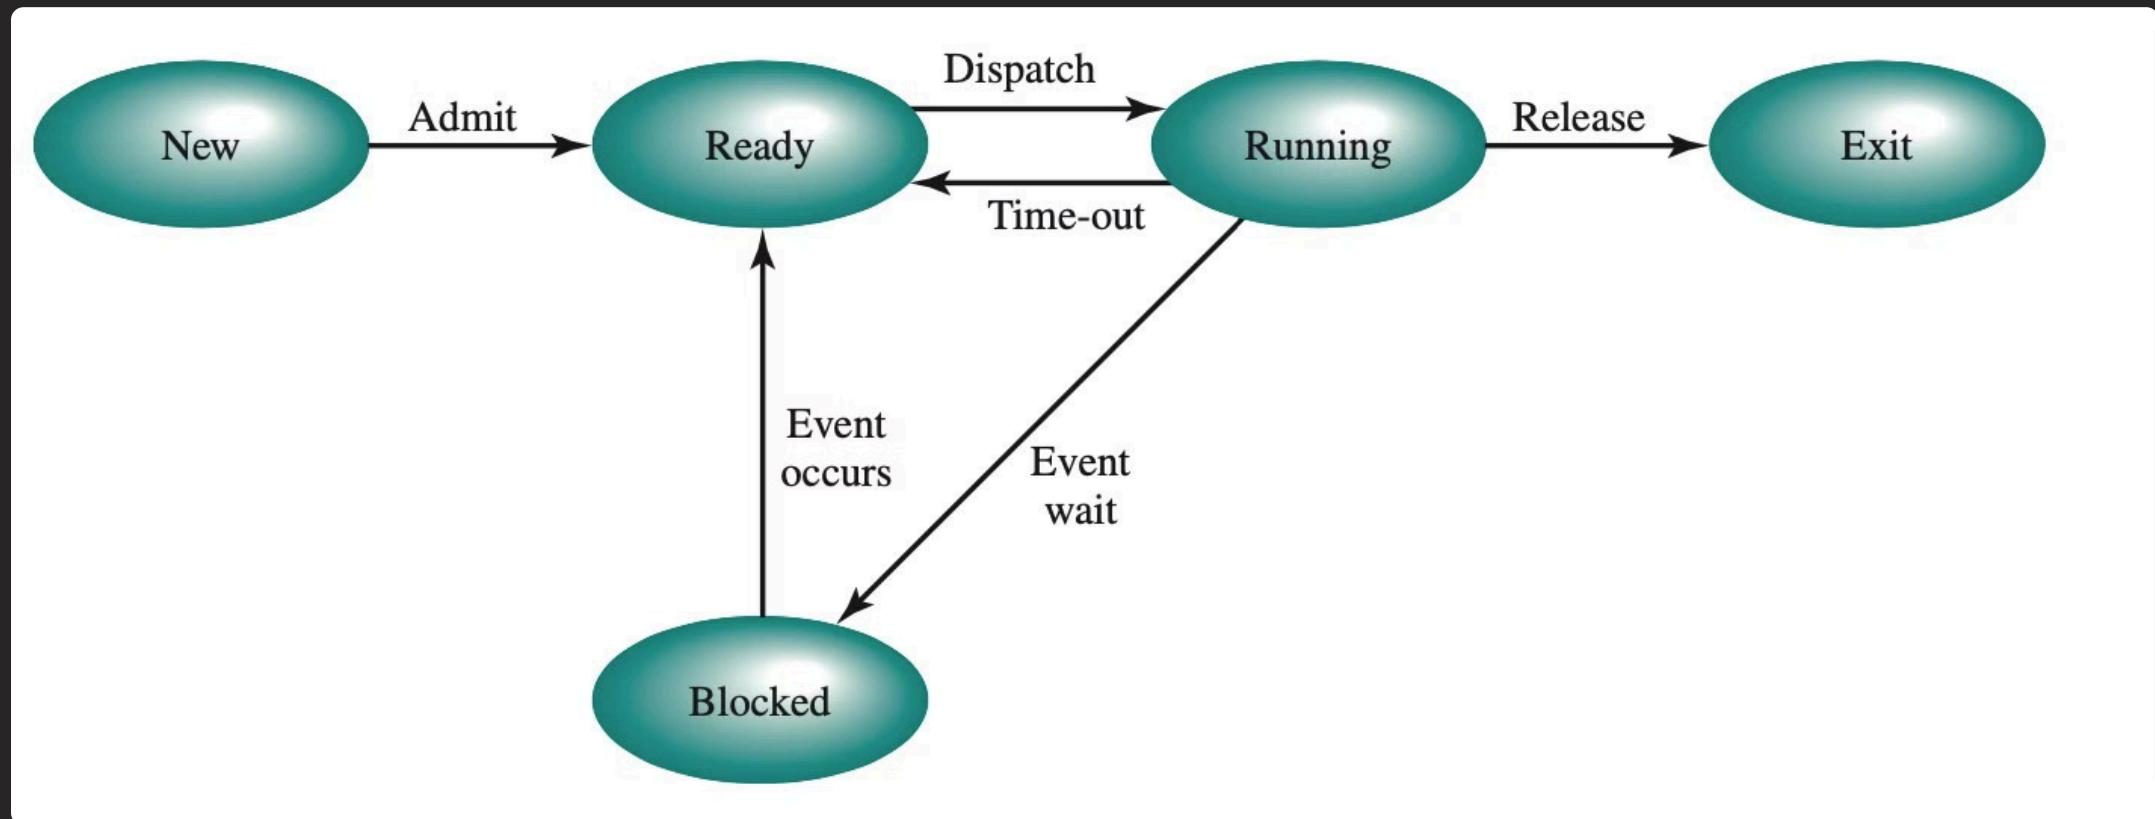
# Why Concurrent Execution?

Efficient resource utilization.

**Fairness:** Equal resource sharing.

**Convenience:** Easier task management.

# Process States



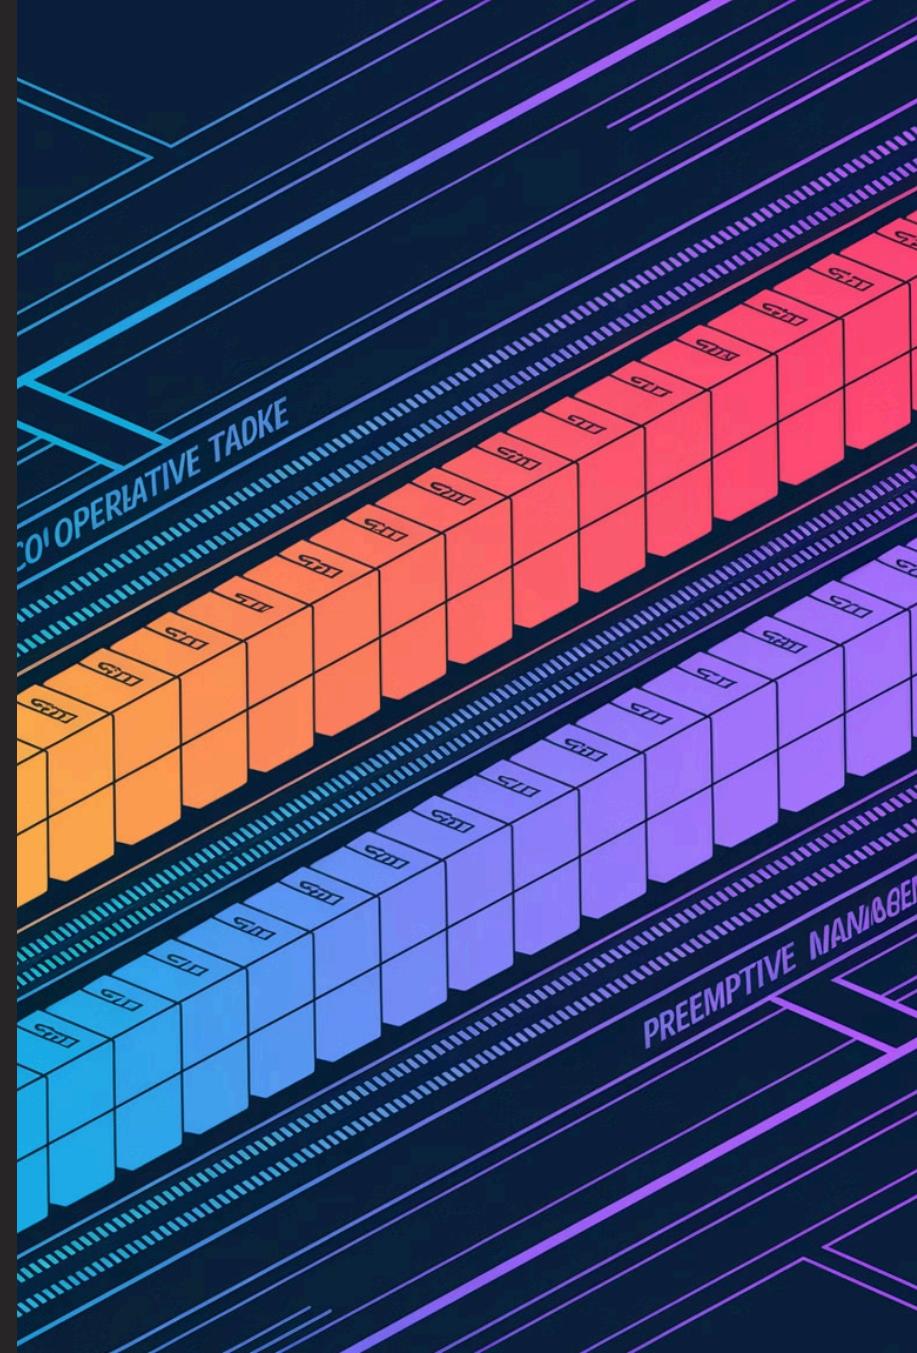
# Scheduling

## Cooperative

Tasks voluntarily yield control of the CPU.

## Preemptive

OS controls task execution, forcing interruptions and resuming them as needed





# Cooperative Scheduling

## 1 Task Control

Tasks control relinquishing the CPU.

## 2 Yielding

Tasks yield when idle or decide to allow others.

## 3 Advantages

Simplicity, **low overhead**, predictable.

## 4 Challenges

Vulnerable to misbehaving processes, potential for CPU monopolization.



# Preemptive Scheduling

1

## Controlled by OS

OS determines when a task relinquishes CPU.

2

## Time Slicing

Tasks have CPU time slices.

3

## Advantages

Improved responsiveness and fairness.

4

## Challenges

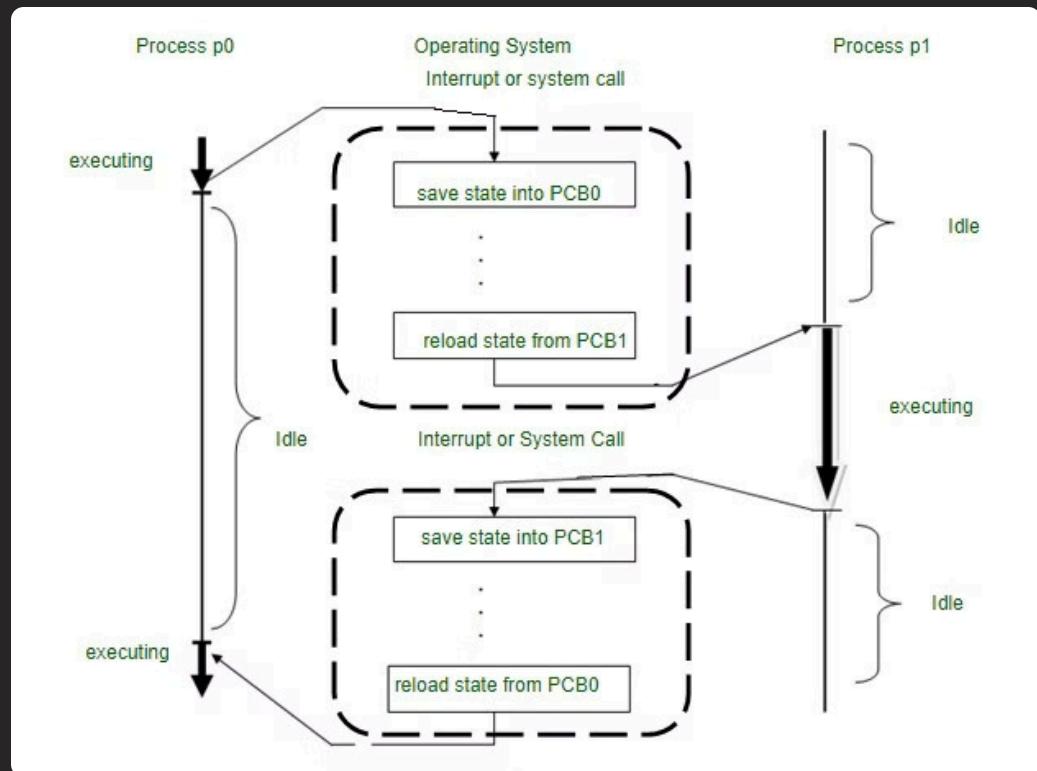
Potential for resource contention, overhead from frequent context switching, and complexity in implementation.



# Context Switch

A context switch stores and restores the CPU state. It pauses and resumes processes.

# Context Switch Process



A context switch involves saving and restoring the execution state when switching between processes.

## Program Counter

Address of the next instruction.

## CPU Registers

Values in processor registers.

## Stack

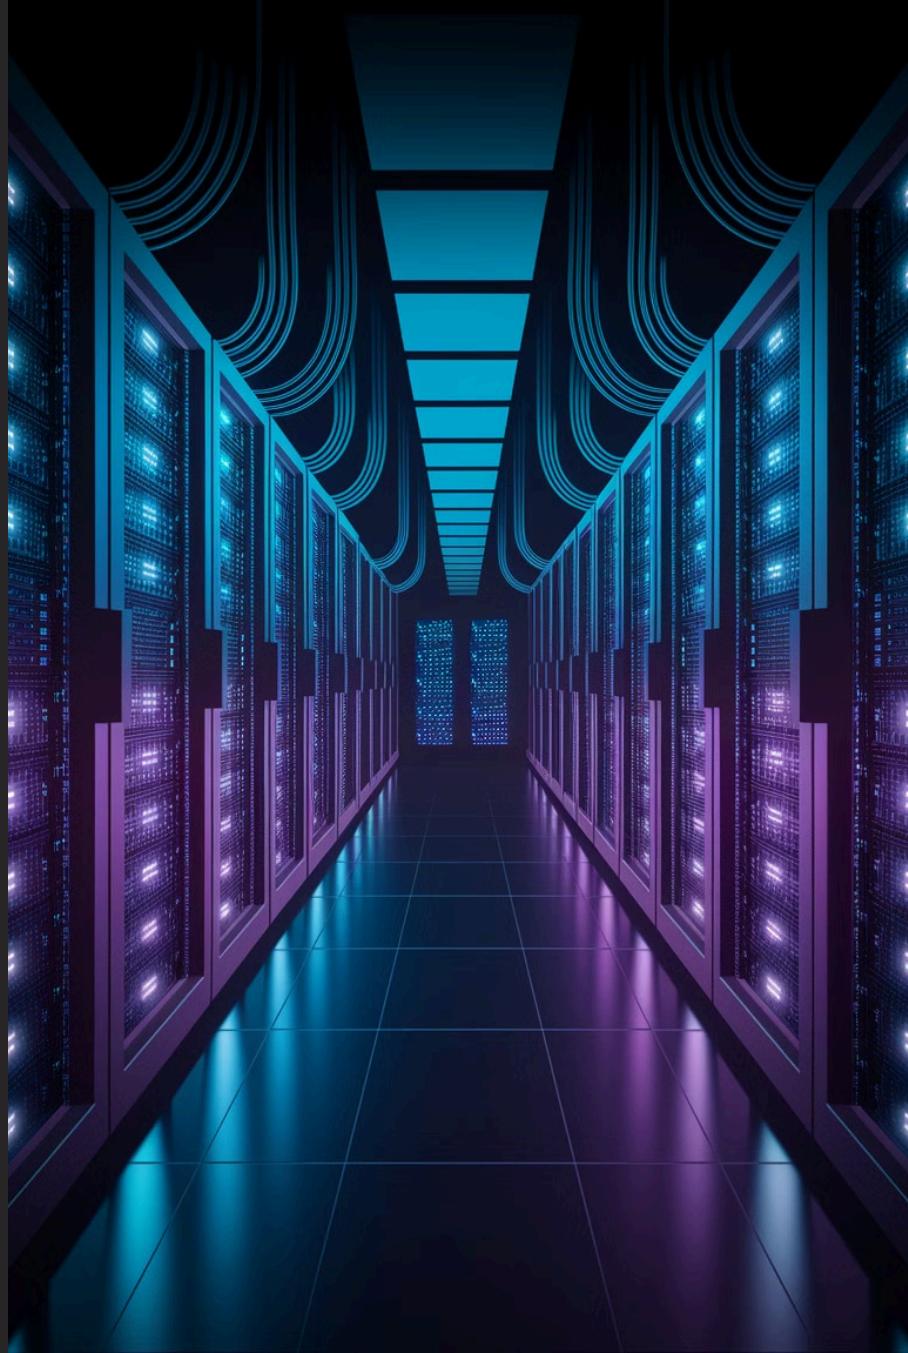
Local variables, parameters.

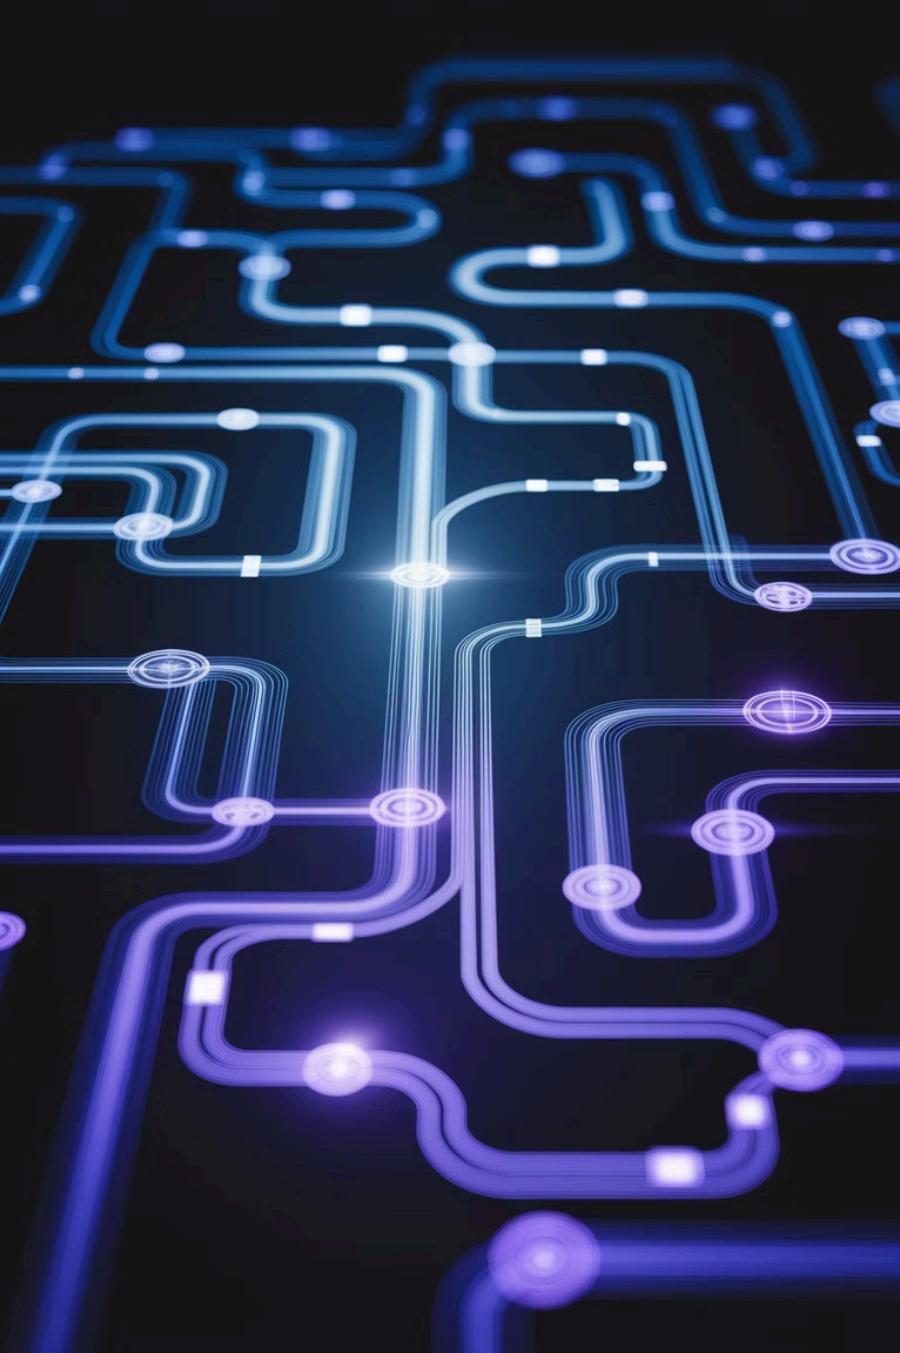
## Memory

Management Information.

# Concurrency vs Parallelism

Concurrency manages multiple tasks. Parallelism executes multiple tasks simultaneously. Interleaving alternates between tasks.

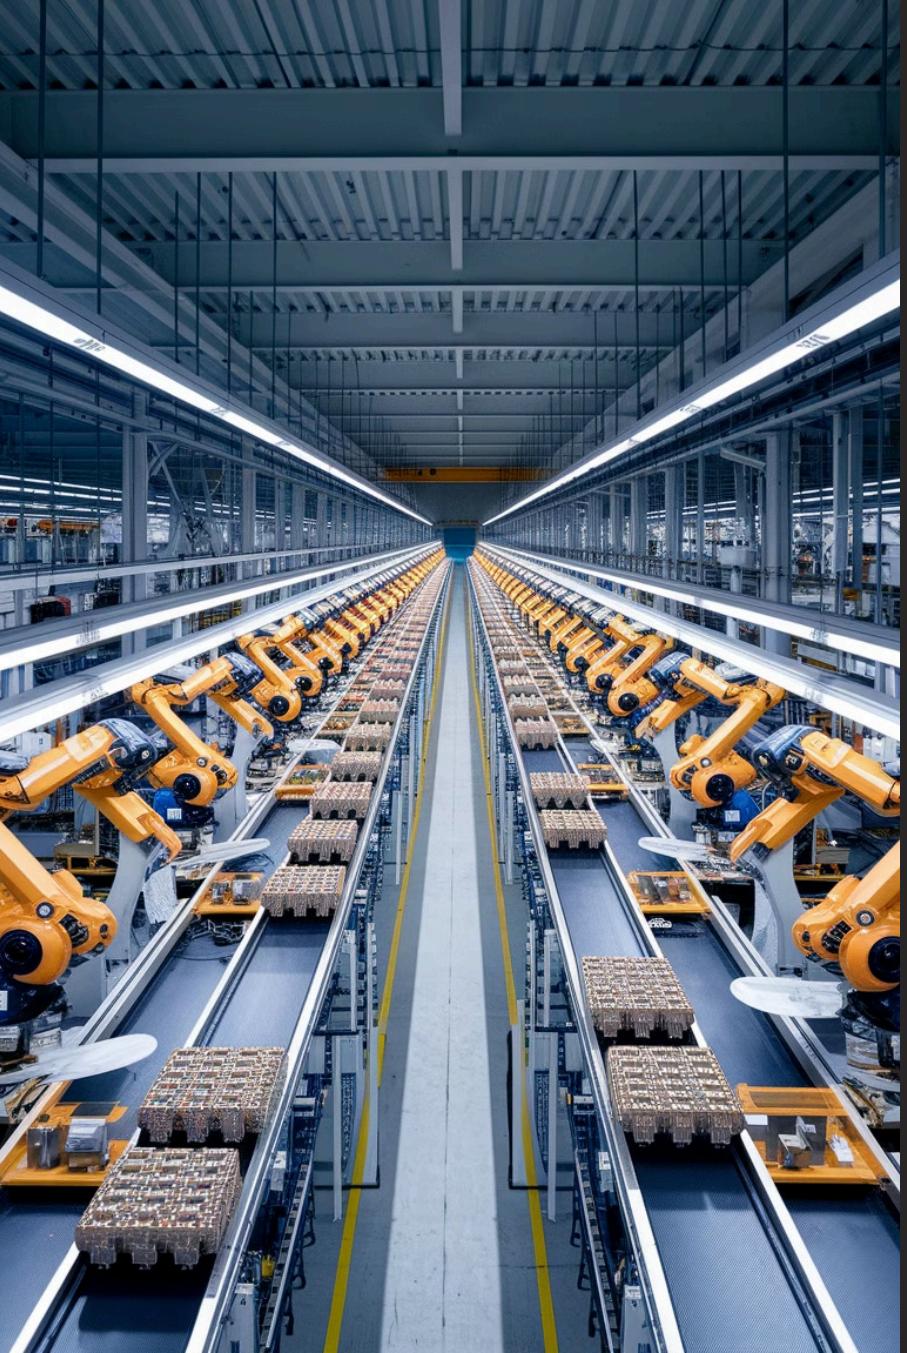




# Concurrency Defined

- 1** The ability of a system to Manage multiple tasks at the same time.
- 2** Operations are not necessarily executing at the same time.
- 3** Different tasks can be in progress at the same time. No necessarily at the same moment.

**Concurrency is about structure**



# Parallelism Defined

- 1 Executing multiple tasks simultaneously.
- 2 Requires hardware with **multiple** processing units.
- 3 Tasks are literally executed at the same time.

**Parallelism is about execution**



# Interleaving

A technique to alternate between tasks rapidly. It executes them in small slices.



# Interleaving Characteristics

- 1 Simulates Parallelism
- 2 Time-sharing  
CPU time is divided among multiple processes.
- 3 Efficiency

Provides efficient CPU utilization.

A chef preparing multiple dishes by switching between them.

Rather than cooking each one start to finish.

# Interleaving and Overlapping

Time →



(a) Interleaving (multiprogramming; one processor)



(b) Interleaving and overlapping (multiprocessing; two processors)

■ Blocked

■ Running

# Processes vs Threads

## Program

A collection of instructions and data that can be executed by a computer to perform a specific task or function.

## Process

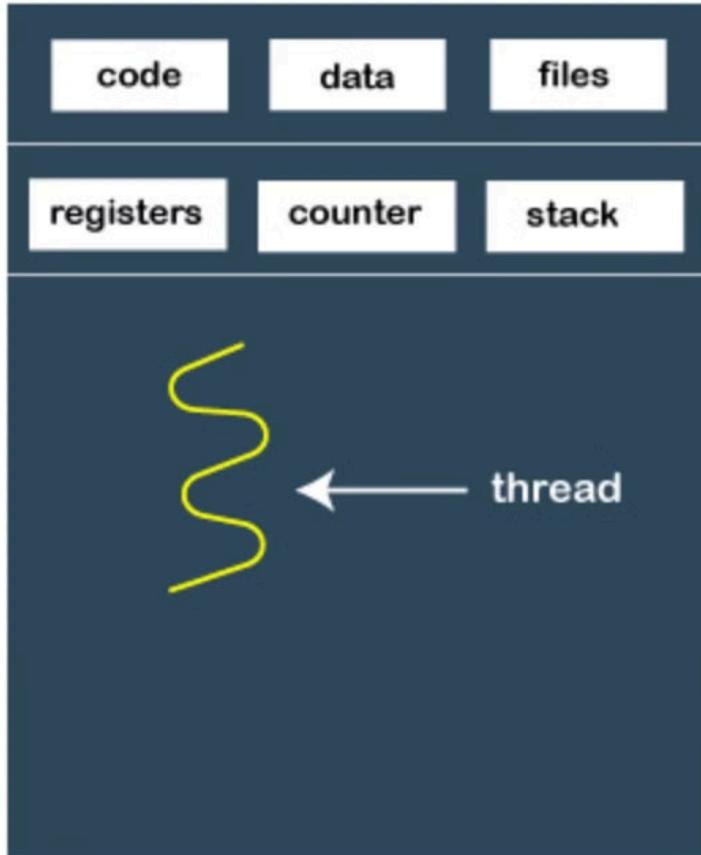
An independent running instance of a program with its own address space, memory, and data.

# Thread Defined

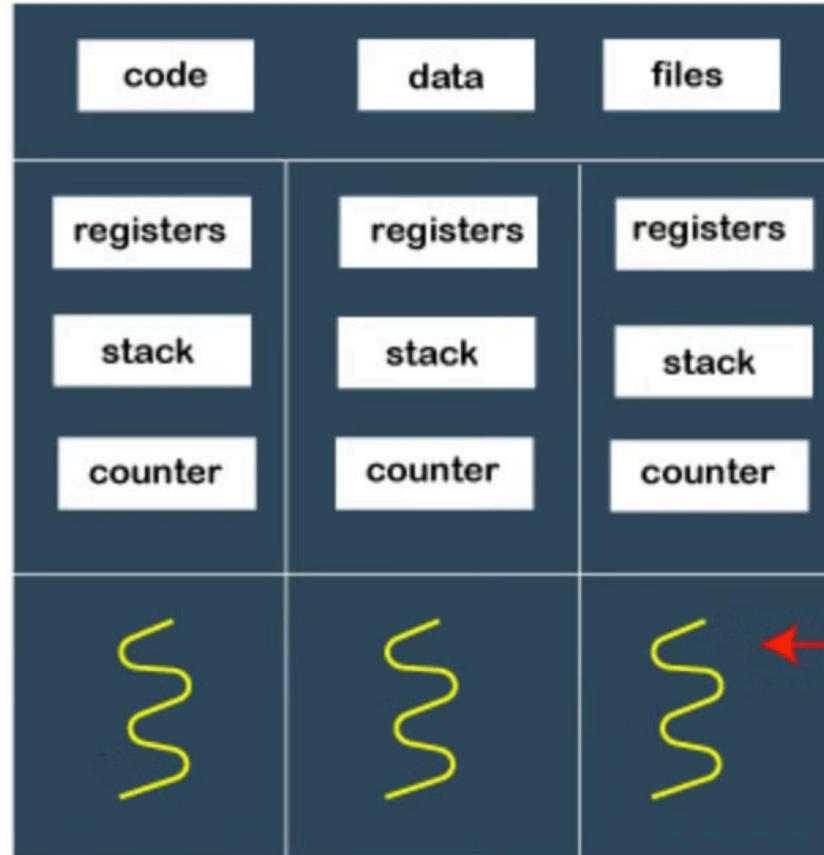
- 1 Smaller unit of execution within a process.
- 2 A single sequence of instructions, that can be scheduled by the system's scheduler
- 3 Multiple threads within a process **share** the same memory space.  
(Not the stack!)



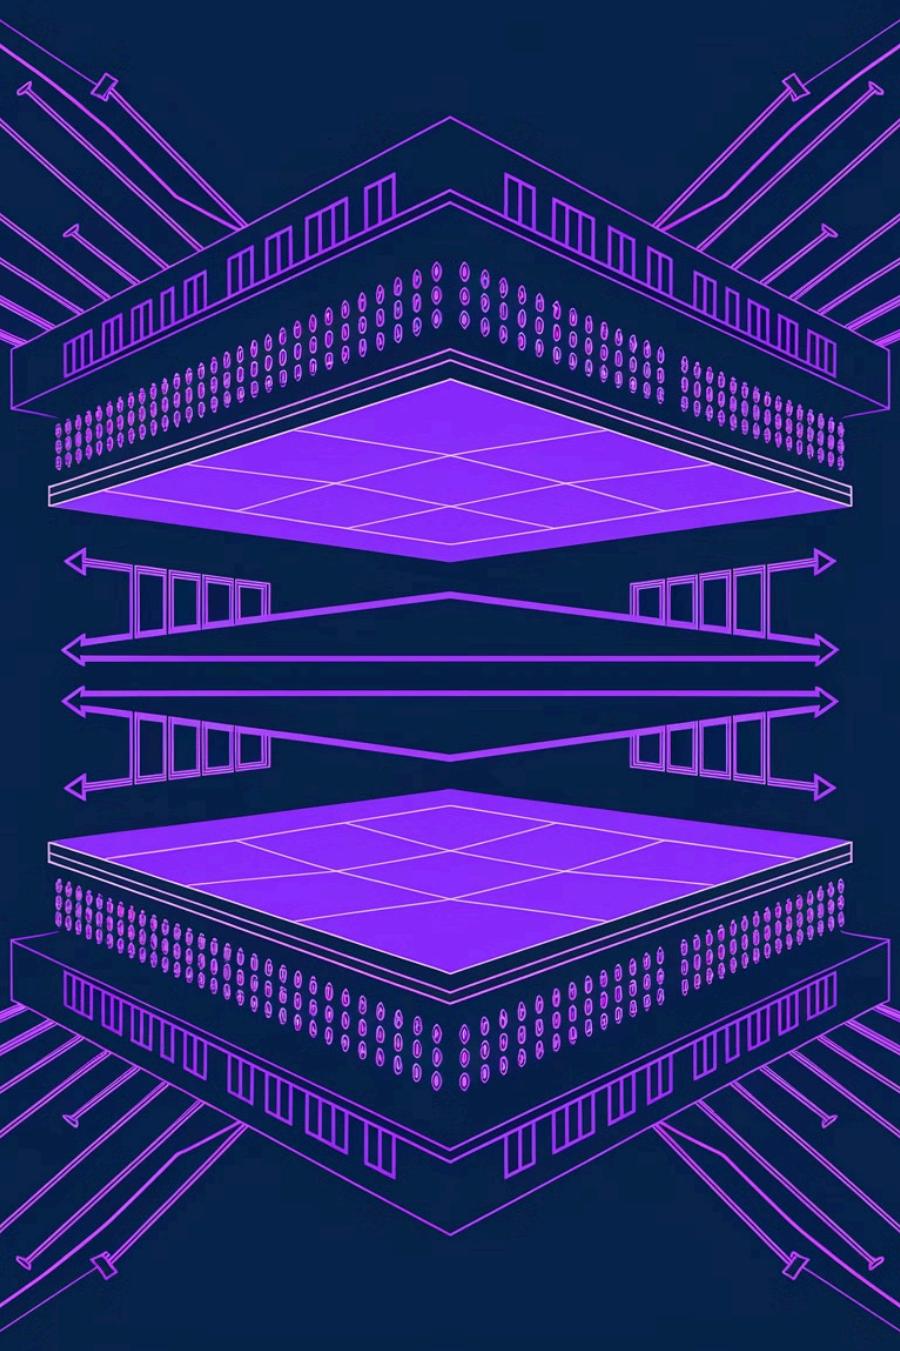
# Processes vs Threads



Single-threaded process



Multi-threaded process



# Types of Threads

Threads can be kernel-level or user-level. Each has different characteristics.



# User-Level Threads

Thread management is done by the application. The kernel is not aware of threads.

**Also called virtual or green threads**



# Kernel-Level Threads

All Thread management is done by the kernel. The application uses an API to interact to the kernel

# Kernel vs User Level Threads

Feature	User Level Threads	Kernel-Level Threads
Implementation	In User Space	In Kernel Space
Context Switch Time	Fast 😊	Slow
Memory Consumption	Low 😊	Higher
Scheduling	By User-Library	By OS 😊
Blocking	One Thread Can Block All	Independent 😊
Multi-CPU Use	None or Limited	Full 😊

# Summary

**1** Concurrent programming history.

**2** Scheduling types.

**3** Processes and context switch.

**4** Concurrency vs Parallelism.

**5** Threads. Kernel vs user level threads.